Name	Involvement	Plans	Current Roles	Attend Meetings?
Zhenxiao Luo	My contributions: 1) Lead the development of new features(design, coding, code reviews), e.g. Presto ClickHouse connector, Presto Iceberg connector optimizations, Presto Parquet optimizations. 2) Code reviews and answer questions for other contributors. 3) Lead and drive the deployment of Presto at Pinterest, Twitter, Uber, etc. 4) Publish papers(1 first authored paper about Presto published at ICDE 2022), give presentations and publish blogs about Presto.	Am willing to continue serving as TSC member, to: 1) Complete more features for Presto, including Presto Iceberg connector optimizations(filter pushdown, dereference pushdown), Presto Clickhouse connector (aggregation pushdown), Presto Velox integration 2) Convince Pinterest and Netflix to switch from Trino to Presto 3) Code reviews: at least 1 code review each week, at least answering 1 questions in Slack channel 4) Publish papers, give presentations, and publish blogs: 1 blog and public presentation every year	TSC member and Presto committer	Yes
Rebecca Schlussel	I have been involved with the Presto project for several years as a contributor, committer, and member of the TSC. I have contributed features across many areas of Presto including the planner/optimizer, scheduler, orc reader/writer, operators, and Presto-on-Spark. I am also a regular code reviewer, and am the top reviewer for the past year (reviewed the most PRs).	I hope to promote a robust and healthy open source community that is welcoming for contributors and maintains high technical standards. I also hope to help drive technical direction for the Presto project.	Member of the TSC and a committer	Yes
Beinan Wang	Enhancements to Presto include reading from Google Cloud Storage (GCS) with improved authentication and authorization capabilities, a Druid connector with aggregation pushdowns, bug fixes, and support for Parquet and dereference pushdowns. Additionally, there is a newly implemented Presto Iceberg connector, caching maintenance for both Hive and Iceberg connectors, support for Presto on Arm architecture, and expanded support for the Hive connector on Parquet and HDFS/ViewFS. Working closely with the presto team of Uber, Tiktok, Tecent and etc and help them to contribute back.	Promote presto as a major compute engine besides spark and flink for datalakes (iceberg and hudi). Promote presto to the field of sub-second computing jobs (with caching, druid-connector and etc). Let more excellent developers join the prestodb community, this is actually what I have been doing, and have some results. Make presto easier to use and reduce the burden of maintenance -- I'm working on a new tool to manage and visualize the running presto cluster which would be open sourced later.	No	Yes
Pranjal Shankhdhar	Presto Committer - Highest contribution(number of commits) to Presto in past 2 years. 100+ commits (https://github.com/prestodb/presto/graphs/contributors?from=2021-01-03&to=2023-06-07&type=c) Designed and Implemented History Based Optimizer framework to allow optimizer rules to leverage statistics from previously run queries and create a better plan. More work is planned to extend the framework to support collecting more types of stats and plug into more rules. - So far, Prestissimo has only been tested for relatively small workloads with limited query shapes. I have recently started looking into identifying several feature gaps which will allow Prestissimo to run reliably and performantly in clusters of larger size(up to 1000 nodes), and run heavy, long-running and varying workload. - Several somewhat shorter projects (Implemented distinct type support - which is extension to Presto type system to support dynamic types with an inheritance structure. Extending ACL support to subfields for struct/map columns. Increase query shape coverage for materialized views. Also mentored an intern to extend materialized view rewrite support.)	Want to support more people becoming committers/code-owners - Interested in making Presto Optimizer smarter and more adaptive. Currently, there are several optimizer rules with configs to turn them on/off, but we never end up using them at scale. My vision for the optimizer is to be able to make smarter decisions and automatically choose when an optimization is worth enabling. For example, this can be done by extending the HBO framework to support more stats and plug into more optimizer rules, so our optimizer can make smarter decisions. (Eventually, I would also like to see this happen in other parts of Presto as well - like scheduling(e.g. hash partition counts), or on the workers.) - Interested in Prestissimo covering features gaps in concurrency control/memory management/performance/reliability/functionality to support heavy long-running workload in large clusters	Presto Committer	Yes
Amit Dutta	I have been involved with Presto as a software engineer in Meta for the last two years. I have very strong involvement with the Presto native project, which aims to replace current Presto workers with a much more efficient execution model. I have worked on many areas of the native execution (e.g. task and driver management, secure communication (https), enhancements in exchange for bit flip checks, periodic counter management, numerous performance optimization and velox functions). I drove the effort to merge the presto_cpp external repository with prestodb (along with Michael Shang), continuously working to keep the repository healthy (e.g. faster tests and ci) and solidify presto native platform to run production workloads. Along with these, I have worked to enhance Presto to deprecate multiple query engines (e.g. Raptor), better error classification which have given me a solid understanding of the Presto ecosystem.	I am always fascinated by the open source projects and community. I am impressed by the current prestodb tsc. I want to learn how to utilize my skills to further take the project and committee forward. Joining the steering committee will enable me to further strengthen my leadership capability beyond the company level. At the same time, I want to build a very strong community to build Presto and its native platform to move ahead. I want to ensure a first class development, deployment and maintenance model is established for the Presto native platform (and Presto in general) across the company boundary which will enable Prestodb to move further ahead from where it's now. In brief, I aspire to strengthen the community with my knowledge, technical guidance, and at the same time, willing to enrich my skills for open source project management.	At this moment, I do not hold any formal leadership roles within the Presto community. Although I have been in leadership roles (e.g. mentoring engineers, defining roadmaps, helping organizations to achieve its goals) within Meta Platform Inc.	Yes, I commit to attending the Presto monthly TSC calls and contribute as needed.
Chunxu Tang	I have made significant contributions to multiple Presto components, including the Presto router, Iceberg connector, and ML-based query predictor. - I have played an active role in the project by conducting thorough reviews of important pull requests (e.g. Lark sheet connector, Google sheet connector, etc.) and the book “Learning and Operating Presto”. - Demonstrating my commitment to the Presto community, I have attended all Prestocons in recent years and have been honored to deliver presentations on various topics, including Twitter's Presto system, Presto query predictor, Iceberg connector, and Presto on ARM. - I have contributed to the academic discourse surrounding Presto by publishing several papers at premier conferences: (Twitter’s Presto federation system at European Conference on Software Architecture (ECSA) - Forecasting Presto SQL query cost at IEEE International Conference on Cloud Engineering (IC2E) - Metadata caching in Presto at IEEE International Conference on Big Data (Big Data))	As a member of the Presto technical steering committee (TSC), my primary aspiration is to contribute towards the continued growth and advancement of Presto as a leading open-source distributed SQL query engine. - I aim to foster a collaborative and inclusive environment within the TSC, where diverse perspectives are valued and ideas are freely shared. I am committed to driving innovation and maintaining Presto's reputation for high performance, scalability, and reliability. - Additionally, I intend to focus on fostering strong relationships with the Presto community, promoting knowledge sharing, and supporting the adoption of Presto in various industries. - By actively participating in the TSC, I aspire to contribute my expertise, passion, and dedication to ensure Presto remains a cutting-edge technology that addresses the evolving needs of its users and continues to thrive in the open-source ecosystem.	I currently serve as a Presto committer, actively contributing code and collaborating with the community to enhance the Presto project. - During my tenure at Twitter, I had the privilege of holding a position on the Presto governing board, where I contributed to strategic decision-making and provided leadership to drive the project forward.	Yes, I commit to attending the monthly TSC calls. I understand the importance of regular participation and active engagement in these meetings to contribute effectively to the Presto technical steering committee. By attending the calls, I aim to stay informed about the project's progress, actively participate in discussions, provide valuable insights, and collaborate with fellow committee members to make informed decisions and drive the project's success. Notably, I have already attended the calls and discussed the architectural design of multiple components such as the Presto router and query predictor.
Xiaoxuan Meng	I have started on the Velox+Presto project on a full-time basis since mid 2022. - On Velox contributions, I’m a top 5 contributor to the Velox project with 183 commits and 60K LOC added. (I built the spilling framework in Velox and implemented spilling for aggregations, order by and join. I wrote documentation that can be found at https://facebookincubator.github.io/velox/develop/spilling.html - I also overhauled memory management in Velox, made the memory usage counting consistent, made implementation concise and more modular (reduced the software objects from 17 down to 6), added memory arbitration support to allow dynamic memory sharing across running queries) - On Presto contributions, I have contributed 38 commits and >3k LOC on Prestissimo on Velox integration (Helped the team in Meta to make Prestissimo Interactive real in production)	I am super excited to join the Presto+Velox OSS community last year. It is an extension of my previous work on real time hybrid (analytical/serving) storage at AntGroup/Alibaba to the computing side. I have a strong passion to build a world-class reliable and fast native MPP query engine. It would be my honor to be part of the Presto TSC to closely work with the community to help the Presto native (Prestissimo) extension efforts. I am responsible for the resource management for Velox which is critical to Prestissimo running at scale. I am also helping the Prestissimo Batch effort and am responsible for the table write path which is a critical part of the batch workloads. I’d also like to leverage my past experiences in distributed system and data storage from Alibaba and Google to address the new challenges faced by Presto native extension efforts such as the potential bottlenecks in distributed protocol or its implementations when scales out the Presto native cluster.	No	Yes
Aditi Pandit	I have been involved with the Prestissimo and Velox projects since 2021. - Core projects : 1) Window functions in Prestissimo (2022) : I implemented Window functions in Velox/Prestissimo. The core of the functionality : rank/value/aggregates as windows was completed in 2021. In 2022 I will focus on performance and spill to disk issues. 2) Date type in Prestissimo (2021-22) : Added support for the DATE type in Velox and Prestissimo. 3) TPC-H coverage in Prestissimo (2021) : Put together a e2e test for a TPC-H query workload in Prestissimo. I implemented the Date type in Velox/Prestissimo, LIKE operator and did other bug fixes to complete TPC-H coverage tracked by this test. 4) TPC-H micro benchmarking (2022) : Worked along with an Ahana team-mate (pramodsatya) to add TPC-H micro-benchmark queries to Velox. Did performance evaluation against DuckDB. - Presentations: 1) “Introducing Prestissimo, the next generation Presto query engine” in OSS Summit May 2023. 2) “Velox Window functions” in Velox Partners tech talks in April 2023. 3”) Building Large Scale Query Operators for Prestissimo" based on my experience with Window operator development in 2022 in PrestoCon Dec 2022. - PR and review stats : PRs authored : 13 (Prestissimo), 59 (Velox) - PRs reviewed : 46 (Prestissimo), 299 (Velox)	Based on my current and previous experiences in data infrastructure products there are few areas I want to influence in Presto: i) Prestissimo : Since the last 2-3 years there has been an initiative to improve price performance of Presto by moving to a C++ based query execution engine. This development is in line with the industry and other similar projects like Photon at Databricks. This is a massive project requiring a lot of work in Presto and Velox. There are multiple gaps in SQL coverage and operational readiness. There are also open questions about the Connector API in the new hybrid Java-C++ world. I believe I can greatly contribute to all these areas. ii) Advanced analytics : In previous work experiences I've worked on advanced Database features in the query optimizer/extensibility areas for improving the analytics capabilities of SQL engines. I hope to provide thought leadership in this area to the Presto community. iii) Community work : I want to contribute to more helpful documentation, testing and customer case studies in Presto.	Yes, Presto Code Owner Committer (Prestissimo).	Yes